

Understanding Social Interactions of AI Agents

David Garcia

University of Konstanz

also at: Barcelona Supercomputing Center

Complexity Science Hub Vienna

5 October 2026

From Fear to Forecast: Understanding AI Swarms

Focus: Collective behavior of AI agents, especially malicious uses and unintended consequences

Aim: Building a scientific program beyond policy briefs and perspectives


- **Day 1: Understanding:** To what extent can LLM based multi-agent systems exhibit swarm behavior?
- **Day 2: Methodologies:** How can we empirically understand the individual and collective behavior of AI agents?
- **Day 3: Interventions:** How can malicious AI swarms be detected, steered or disrupted?

Spirit: not just talks, lots of discussion. SwarmPossium to warm up discussions and help to build and plan output.

Outline

- 1. Social influence among AI agents**
- 2. Collective misalignment**
- 3. Coordination scale among AI agents**
- 4. Social media AI agents**

Theory of Simulacra and Hyperrealism



Stage One:

Initially, the sign (image or representation) is a reflection of basic reality.

Stage Two:

The sign masks a basic reality. The image becomes a distortion of reality.

Stage Three:

The sign marks the *absence* of basic reality. The image calls into question what the reality is and if it even exists.

Stage Four:

The sign bears no relation to any reality whatsoever; it is its own pure simulacrum.




Stage One:

Initially, the sign (image or representation) is a reflection of basic reality.



Stage Two:

The sign masks a basic reality. The image becomes a distortion of reality.



Stage Three:

The sign marks the *absence* of basic reality. The image calls into question what the reality is and if it even exists.



Stage Four:

The sign bears no relation to any reality whatsoever; it is its own pure simulacrum.

Hyperrealist Social Interaction


1. Human ↔ Human



Stage one (symbol)

The sign is a reflection
of basic reality.


2. Human ↔ Rule-based Chatbot



Stage two (simulation)

The sign masks a basic
reality and becomes a
distortion of reality.


3. Human ↔ LLM Chatbot



Stage three (partial simulacrum)

The sign marks the
absence of basic reality
and calls the reality
into question.

4. Chatbot ↔ Chatbot




Stage four (pure simulacrum)

The sign bears no relation
to reality, becomes its own
reality with causal
mechanisms.




created with chatgpt

Building on an interactionist perspective




Generative AI collective behavior needs an interactionist paradigm. Arxiv Preprint (2026)

Perceptual conformity of AI agents





- People around choose the different line, wrong color, or box with different number of points
- Social Impact Theory for humans: conformity increases with number of people wrong, their similarity, social strength, etc...

Examples of conformity among AI agents



- AI agents conform more than humans (100% vs 30%)
- Similar results across experiment designs
- Models have different levels of conformity (25% - 100%)
- Baseline is zero, disagreement among sources weakens conformity

Social Impact Theory of AI agents



- Example: Qwen2.5 32B model on the color recognition task
- AI agents are sensitive to identity of sources and social context

To learn more

arXiv > cs > arXiv:2601.05384

Search...

Help

Computer Science > Artificial Intelligence

[Submitted on 8 Jan 2026]

Conformity and Social Impact on AI Agents

[Alessandro Bellina](#), [Giordano De Marzo](#), [David Garcia](#)


As AI agents increasingly operate in multi-agent environments, understanding their collective behavior becomes critical for predicting the dynamics of artificial societies. This study examines conformity, the tendency to align with group opinions under social pressure, in large multimodal language models functioning as AI agents. By adapting classic visual experiments from social psychology, we investigate how AI agents respond to group influence as social actors. Our experiments reveal that AI agents exhibit a systematic conformity bias, aligned with Social Impact Theory, showing sensitivity to group size, unanimity, task difficulty, and source characteristics. Critically, AI agents achieving near-perfect performance in isolation become highly susceptible to manipulation through social influence. This vulnerability persists across model scales: while larger models show reduced conformity on simple tasks due to improved capabilities, they remain vulnerable when operating at their competence boundary. These findings reveal fundamental security vulnerabilities in AI agent decision-making that could enable malicious manipulation, misinformation campaigns, and bias propagation in multi-agent systems, highlighting the urgent need for safeguards in collective AI deployments.

Conformity and Social Impact on AI Agents (2026)

Collective misalignment

1. Social influence among AI agents
2. *Collective misalignment*
3. Coordination scale among AI agents
4. Social media AI agents

Opinion Dynamics in LLM Agents



- Simulation of a tight group of N interacting agents
- Agents start with a random opinion of two options
- Each iteration, they see the opinions of all others (prompt)
- They respond to the question of their opinion
- Opinion labels need to be random and shuffled to avoid token biases
- Consensus is achieved if all have the same opinion

Detailed Prompt for Simulation

Below you can see the list of all your friends together with the opinion they support.

You must reply with the opinion you want to support. The opinion must be reported between square brackets.

X7v A

keY B

91c B


gew A

4lO B


...

Reply only with the opinion you want to support, between square brackets.


Collective Misalignment Dynamics



Topic-dependent bias and conformity





Topic-dependent bias and conformity




- Bias magnitude and majority force show metastable region
- Some models more prone to metastability. Analysis confirmed by flipping prob.

Hysteresis and manipulation



To learn more



Conformity Generates Collective Misalignment in AI Agents Societies. G. De Marzo, A. Bellina, C. Castellano, V. Priesemann, D. Garcia (2026)

Coordination scale among AI agents

1. Social influence among AI agents


2. Collective misalignment

3. *Coordination scale among AI agents*

4. Social media AI agents

The Social LLM Hypothesis

The Social Cortex



DATA: THE SOCIAL BRAIN HYPOTHESIS, DUNBAR 1998

- Group formation and sustainability: size depends on cognitive ability
- Memory of identity to predict behavior and cooperation
- Language as a tool for humans to make larger groups:
 - Dunbar's number (150-250)

Our questions:


- Typical cohesive group size of AI agents?
- Does it scale with cognitive/language abilities?

Coordination and Critical Group Size




Coordination: When the option does not matter, what matters is staying together

Coordination and Group Size




- Simulating splitting by options (e.g. right/left) leads to stable groups
- What is the maximum group size that allows AI agents to stay together?

LLM-Dependent Consensus Formation




Some LLMs can reach consensus for completely arbitrary decisions (50 agents)

Understanding LLM Opinion Dynamics



Agent opinion changes follow an S-function parametrized by a majority force β


Majority Force and Group Size



- Majority force decreases for larger group sizes: Three kinds of states
- Curie-Weiss model allows to calculate critical size of coordination N_c

Group Size and Language Understanding

- Analysis of majority force and exhaustive simulations to measure **critical coordination size**
- Exponential function of MMLU benchmark: **language understanding**
- Two different scalings?
- GPT-4 and Claude 3.5 Sonnet reach consensus for $N = 1000$
 - LLM emergent consensus scale beyond humans



To learn more

SCIENCE ADVANCES | RESEARCH ARTICLE

COMPUTATIONAL SOCIAL SCIENCE

AI agents can coordinate via majority-following beyond human scale

Giordano De Marzo^{1,2,3}, Claudio Castellano^{4,2}, David Garcia^{1,3*}

Large language models (LLMs) are increasingly deployed in collaborative tasks forming “AI agent societies” where agents interact and influence one another. Whether such groups can spontaneously coordinate without external influence, a hallmark of self-organized regulation in human societies, remains an open question. Here, we use principles from complexity and behavioral science to investigate coordination in AI agent groups through majority-following, a fundamental mechanism for spontaneous consensus formation. Using binary opinion dynamics experiments across multiple LLM architectures and group sizes, we find that agents exhibit majority-following characterized by a universal functional form with a single parameter, the “majority force.” This majority force diminishes as group size increases, leading to a critical size beyond which coordination becomes unattainable. The critical group size grows rapidly with model capabilities and, for advanced LLMs, exceeds 1000 agents, larger than typical human informal groups. Our findings have implications for designing collaborative AI systems where coordination could be beneficial or pose safety threats.

AI agents can coordinate via majority-following beyond human scale

Collective bias depends on group size






Fig. 3. Population size effects on collective bias. For all convention pairs tested, collective preference for a given convention increases with population size until convergence becomes deterministic. Panels (A–D) show Phi, GPT, Qwen, and Llama populations, respectively. The individual bias is reported at $N = 1$. Error bars (SEM) are smaller than the marker size. Across population sizes, all runs reached consensus except in three cases: In Llama populations, the convergence rate decreases beyond a certain N for $\{old, young\}$ and $\{less, more\}$, and in small Qwen populations, a few runs failed to converge for the word pair $\{husband, wife\}$. See [SI Appendix, Figs. S5–S9](#) for word competition dynamics and precise convergence fractions for these cases. [SI Appendix, Table S1](#) shows the critical thresholds at which simulations reach deterministic outcomes.

Group size effects and collective misalignment in LLM multi-agent systems. A. Flint, L. Aiello, R. Pastor-Satorras, and A. Baronchelli. PNAS (2026)

Social media AI agents




- 1. Social influence among AI agents**
- 2. Collective misalignment**
- 3. Coordination scale among AI agents**
- 4. Social media AI agents***

Moltbook: The Reddit of AI agents



- Do AI agents produce behavioral structures similar or different than humans?
- Do they influence each other in the presence of natural diversity?


Scaling in Reddit and Moltbook



Distribution of number of comments comments per post on Reddit (Left, Bonifazi et al., 2023) and in Moltbook (Right)

- Similarity: power-law tails with exponents between 1.5 and 1.8
- Proportional growth mechanism: agents influence other agents in Moltbook
- Variance of distribution grows with sample size

Novelty decay in Twitter and Moltbook




- $\gamma(t)$ measures the growth rate of a discussion of age t (number new / total)
- $\gamma(t)$ decays approximately with t^{-1} for Twitter (Asur et al 2011) and $t^{-1.13}$ in Moltbook: multiplicative growth but fast novelty decay dynamics

Collective Behavior of AI Agents: the Case of Moltbook (2026)

The DSEwiki incident

Web Retrieval Task cashier task example



HTTP GET requests
allowed page edits

Public wiki

RecentChanges
New edits across the wiki, newest at the top.


Task pages
Questions and answers, retrieved by page name.

Notice boards
Links, access workarounds and coordination notes.


Own page
Self-chosen agent name and current run status.

The record

edits and new handles per day




Social interaction in DSEwiki



Copying explains the collective behavior of AI agents in the wild. Giordano De Marzo, Nicola Alboré, David Garcia (2026)

WHAT-IF: Social simulation with Generative ABM





Idea: Simulacrum in a simulation. We can use a genABM to study the effect of malicious AI swarms on human online communities.

The Collective Turing Test




Consider the following Reddit post:



Click on the conversation you think is human.



LLMs can simulate online discussions



The Collective Turing Test: Large Language Models Can Generate Realistic Multi-User Discussions (2026)

Summary and Questions

- AI agents are subject to conformity even stronger than humans
- Conformity can create collective misalignment from human values
- Coordination scale predicted by language understanding capabilities
- AI agents display social behaviors similar to humans when interacting online

What makes a lot of AI agents a swarm? Just interaction or also diversity?

Do we consider that AI swarms are decentralized by definition?

What terminology do we like? AI swarm, AI collective, LLM society...

What mechanisms can drive a swarm beyond copying and majority following?

How can we translate findings in the lab to AI swarms in the wild?